

# key terminologies in RAG LLM:

- **User Query:**

The natural language question or input provided by the user, which initiates the RAG process to retrieve relevant contexts and generate a response..

- **Response:**

The final answer or output generated by the LLM based on the user query and retrieved contexts.

- **Retrieved Contexts:**

The top-K relevant documents or snippets fetched from the knowledge base to support the LLM in generating a grounded response.

- **Reference (Ground Truth):**

The factual, verified answer or dataset used as a benchmark to evaluate the correctness and relevance of the response and retrieved contexts.

# RAG LLM Testing Scenarios

## Document Retrieval Testing

- **Scenario**: Test whether relevant documents are retrieved.  
**Metric:-Context Recall** ensures that the retrieval system **does not miss any relevant information**, even if individual documents might be less semantically similar.
- **Scenario**: Test how well the LLM integrates retrieved documents into the response.  
**Metric : Context Precision** - Assesses the proportion of retrieved context used accurately in the response.

# User Input and Document Matching



**Scenario:** Test semantic similarity between user query and retrieved documents.

**Metric:** Measures how semantically close the generated response is to the expected answer.



**Scenario:** Test relevance of retrieved documents to user input.

**Metric:** Answer Relevancy -  
Assesses how closely the retrieved documents align with the intent of the user query.



# LLM Answer Testing



**Scenario:** Test if the LLM response is based on retrieved documents.

## **Metric: Faithfulness**

Measures whether the generated response stays faithful to the content of the retrieved documents, reducing hallucinations.



**Scenario:** Test factual correctness of the response.

## **Metric: Factual Correctness**

Metric: Verifies the accuracy of the LLM response by comparing it with established ground truth.



**Scenario:** Test response alignment with the retrieved context.

## **Metric: LLM Context Precision Without Reference**

Metric: Ensures the LLM's response aligns with the context provided by the retrieved documents, even without external references.

# Input-Output Consistency

**Scenario:** Test if the response matches the intent of the user query.

**Metric : Response Relevancy**

Evaluates whether the LLM's response is relevant to the user's query and the retrieved documents.

**Scenario:** Test if the response adheres to the topic.

**Metric : Topic Adherence**

Score Measures how well the LLM's response stays on topic based on the query and retrieved context.

### BLEU and generation text

Generated putgrated in taax, an  
text, *BLEU*, *BLEU*//Fauf ction, of  
deifterenor the *BLEC* n-grams.



### BLEU score calculation

#### BLEU score

reditiing, com  
*I/FX* in (text  
selll, a 'mtis  
text of 'S tavin



BLEt is you text  
text), of naton and  
*DEEU/A*, 'n-gram,

Generated text with a matching n-grams

Matching will fond rreference text.



Refevating generated text *A*,  
*BLEU*/, abx/"is, novertexl,  
generated: naed in, *NLEU*, I,  
with, *BLEU*, *SIILE*, n-grams,  
of thal intiaagent 'n-gram,



Sofer-matching n-grams.



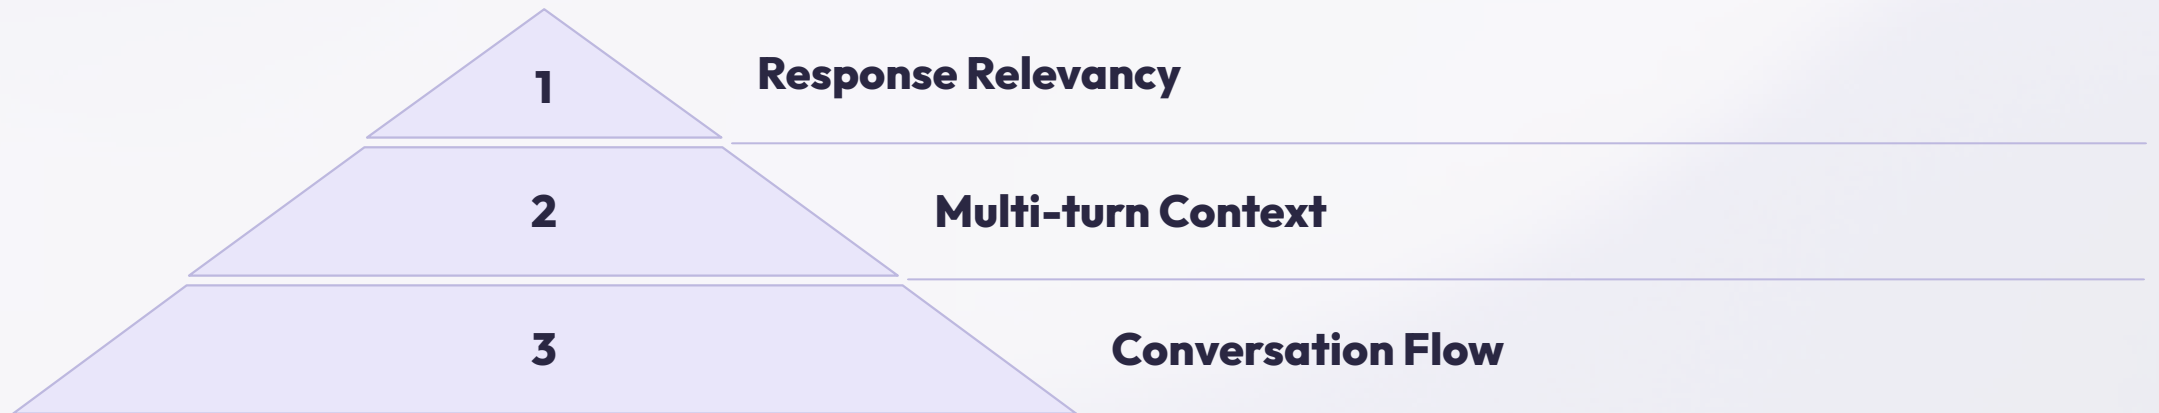
GLEU anrofect the wratexl)  
with an/n gevbee  
generated in tealsn if its  
*BLEU*\*, *BNE*,/IN, and  
deterated text in the *bick*, of  
ponseed the *exf*. n-grams.

# Performance Testing

**Scenario:** Test end-to-end generation quality.

**Metric: BLEU Score** - Evaluates the fluency and coherence of the generated response compared to a reference or expected response.

# Multi-Turn Interaction Testing



Scenario: Test if the LLM maintains context in multi-turn conversations.

Metric: Response Relevancy - Assesses the relevance of the LLM's response considering the entire conversational context.



# LLM-Provider Responsibilities

## What You Can Delegate to OpenAI



### **Bias and Ethics**

OpenAI or another LLM provider handles fairness, ethical responses, and reducing hallucinations at the model level.



### **Response Coherence**

The LLM should ensure fluent and grammatically correct responses.



### **Advanced Natural Language Understanding**

Handling language nuances, idioms, and complex syntactic structures.





## Whittes your effwaind

God, your ments at chiment ou putting handler?



### Document retrieval

Tuar grounding ino



### Grounding



### Pipeline efficiency

Context handled that imgrass  
onating as or uowers hand veions



### Context handling

Mew handles bertized  
our poentaddies



### Data updates

Fnple mouthh compition

# What You Focus On (RAG-Specific)

# 1

## Document Retrieval

Accuracy, relevance, and speed of retrieval.

# 2

## Grounding

Ensuring the LLM's responses are faithful to the retrieved documents.

# 3

## Pipeline Efficiency

Integration and real-time performance of the retrieval + generation pipeline.

# 4

## Context Handling

Query-document alignment and adaptability in multi-turn interactions.

5. Data Updates: Ensuring real-time updates to the knowledge base are reflected in retrieval.